

Lesson 9 -> Partitioning vs Sharding

1 Partitioning vs Sharding (Clear Distinction)

Partitioning

- Logical division of data
- Usually within the same database system
- May or may not be on different machines

| Think: "How is data split inside a DB?"

Sharding

- Physical distribution of data across multiple databases / nodes
- Always implies horizontal partitioning
- Used when a single DB cannot handle scale

| Think: "Which database node holds this data?"

Simple Analogy

- Partitioning → rooms inside one building
- Sharding → multiple buildings

2 Database Partitioning (Techniques)

2.1 Horizontal Partitioning (Row-based)

What it is

- Rows are split across partitions
- Same schema in each partition

Example:

```
orders_2024 orders_2025
```

Each partition contains different rows.

Pros

- Scales reads & writes
- Queries scan less data

Cons

- Cross-partition joins expensive
 - Partition management complexity
-

2.2 Vertical Partitioning (Column-based)

What it is

- Columns split into separate tables
- Based on access patterns

Example:

```
users_basic (id, name) users_profile (id, address, preferences)
```

Pros

- Faster reads for common queries
- Smaller row size

Cons

- Joins needed for full data
 - Schema complexity
-

2.3 Range-Based Partitioning

What it is

- Data partitioned based on value ranges

Example:

```
orders: Jan–Mar → P1 Apr–Jun → P2
```

Pros

- Very efficient range queries
- Good for time-series data

Cons

- ✗ Hot partitions
- ✗ Manual rebalancing

Used heavily in:

- Retail orders
 - Logs
 - ML features by time
-

2.4 Hash-Based Partitioning

What it is

- Hash function decides partition

Example:

```
hash(user_id) % 4 → partition
```

Pros

- ✓ Even data distribution
- ✓ Avoids hot partitions

Cons

- ✗ Poor range queries
 - ✗ Rehashing pain when adding partitions
-

2.5 Key-Based Partitioning

What it is

- Partitioning based on a natural key

Example:

```
user_id → partition
```

Often implemented via hashing.

Pros

- Predictable routing
- Simple logic

Cons

- Key skew causes imbalance
-

2.6 Composite Partitioning

What it is

- Combine multiple strategies

Example:

- Range by date
- Hash by user_id inside range

Used in **high-scale retail & ML systems.**

3 Database Sharding (Techniques)

Sharding is **horizontal partitioning + distribution.**

3.1 Range-Based Sharding

How it works

Shard 1 → user_id 1–1M Shard 2 → user_id 1M–2M

Pros

- Simple routing
- Range queries efficient

Cons

- Hot shards
 - Rebalancing difficult
-

3.2 Hash-Based Sharding

How it works

```
hash(user_id) % N → shard
```

Pros

- Uniform load
- Easy scale (initially)

Cons

- Resharding is painful
 - Range queries inefficient
-

3.3 Consistent Hashing (Important)

How it works

- Hash ring
- Minimal key movement when adding/removing shards

Pros

- Scales well
- Minimal data reshuffle

Cons

- Complex to implement
- Debugging harder

Used in:

- Cassandra
 - DynamoDB
 - Redis Cluster
-

3.4 Directory-Based Sharding

How it works

- Lookup service maps key → shard

Example:

```
user_id → shard_3
```

Pros

- Flexible
- Easy rebalancing

Cons

- Directory is SPOF
 - Extra network hop
-

3.5 Geo-Based Sharding

How it works

- Shards based on geography

Example:

- India → shard IN
- US → shard US

Pros

- Low latency
- Compliance friendly

Cons

- Cross-region queries hard

Used by:

- Global retail platforms
 - Recommendation systems
-

4 Partitioning vs Sharding (Side-by-Side)

Aspect	Partitioning	Sharding
Scope	Logical	Physical

Aspect	Partitioning	Sharding
Machines	Same or multiple	Multiple
Scaling	Limited	High
Complexity	Medium	High
Failure impact	Local	Larger

5 ML & Retail-Specific Examples

Feature Store

- Partition by date
- Shard by entity_id

User Embeddings

- Hash-based sharding
- Consistent hashing

Online Inference Cache

- Redis cluster with hash slots

Training Data

- Range partitioned by time
- Sharded by region

6 Failure & Edge Cases (Architect Thinking)

Hot Partitions

- Skewed keys
- Celebrity users

Solution:

- Salting keys
- Adaptive sharding

Cross-Shard Queries

- Aggregations across shards

Solution:

- Fan-out queries
 - Pre-aggregations
-

Rebalancing

- Adding/removing nodes

Solution:

- Consistent hashing
 - Background migration
-

7 Interview-Ready One-Liners

- Partitioning organizes data inside a database; sharding distributes it across databases.
 - Sharding always implies horizontal partitioning.
 - Hash-based sharding favors uniform load; range-based favors query efficiency.
 - Consistent hashing minimizes data movement during scaling.
-

How This Fits Your Growth Path

Understanding this means:

- You can design **feature stores**
- You can scale **online inference**
- You can discuss **data-heavy system tradeoffs**

This is **ML Architect / Head of DS-level knowledge**.