

Lesson 10 -> Data Replication

1 What Is Data Replication?

Data replication is the process of **copying the same data to multiple nodes**.

So instead of:

Data → 1 database

You have:

Data → DB1 → DB2 → DB3

Each replica stores the **same logical data**.

2 Why Do We Need Data Replication?

1. High Availability

If one node fails, others can serve traffic.

2. Read Scalability

Reads can be distributed across replicas.

3. Low Latency

Serve users from the nearest replica (geo-replication).

4. Fault Tolerance

Protection against hardware failures.

5. Disaster Recovery

Backups + replicas = safety net.

3 Types of Data Replication

3.1 Synchronous Replication

How it works

- Write is successful **only after all replicas confirm**

Client → Leader → Replica1 ✓ → Replica2 ✓ → ACK

Pros

- ✓ Strong consistency
- ✓ No stale reads

Cons

- ✗ High latency
- ✗ Poor availability

Used in:

- Banking
- Financial systems

3.2 Asynchronous Replication

How it works

- Leader commits first
- Replicas catch up later

Client → Leader ✓ → Replica (later)

Pros

- ✓ Low latency
- ✓ High availability

Cons

- ✗ Replication lag
- ✗ Potential data loss

Used in:

- MySQL
 - MongoDB
 - Most web apps
-

3.3 Semi-Synchronous Replication

How it works

- Leader waits for **at least one replica**

Pros

- Balanced consistency
- Lower data loss risk

Cons

- Slower than async
- Complex

Used in:

- MySQL semi-sync mode

4 Replication Models (Very Important)

4.1 Single-Leader (Primary–Replica)

Model

- One leader handles writes
- Replicas handle reads

Writes → Leader → Reads → Replicas

Pros

- Simple
- Strong write consistency

Cons

- Leader bottleneck
- Failover complexity

Real DBs

- MySQL

- PostgreSQL
 - MongoDB (primary-secondary)
-

4.2 Multi-Leader Replication

Model

- Multiple leaders accept writes
- Leaders replicate to each other

Write → Leader A Write → Leader B

Pros

- High write availability
- Geo-distributed writes

Cons

- Conflict resolution
- Complex logic

Real DBs

- Cassandra (conceptually)
 - MySQL multi-master
 - CouchDB
-

4.3 Leaderless Replication

Model

- No single leader
- Clients write to multiple nodes

Client → Node1 Client → Node2 Client → Node3

Pros

- No single point of failure
- High availability

Cons

- ✗ Eventual consistency
- ✗ Read repair complexity

Real DBs

- Cassandra
 - DynamoDB
 - Riak
-

5 Read–Write Strategy

Defines:

- **Where writes go**
 - **Where reads come from**
-

Common Strategies

Write

- Usually to leader (or quorum)

Read

- From leader → strong consistency
 - From replicas → eventual consistency
-

Read Preferences

- Read-after-write consistency
- Read-your-writes
- Monotonic reads

Very important for:

- User-facing ML features
 - Recommendations
 - Personalization
-

6 What Is Replication Lag?

Replication lag = delay between leader write and replica update.

Example:

Write at T₀ → Leader Replica updates at T₀ + 2s

Causes

- Network latency
 - High write throughput
 - Disk IO
 - Replica overload
-

Why Replication Lag is Dangerous

- Users see stale data
 - ML models get outdated features
 - Counters become inconsistent
-

Mitigation

- Read from leader
 - Sticky sessions
 - Version checks
 - Async tolerance
-

7 Database Comparison (Read–Write & Replication Models)

DB	Replication Model	Write Strategy	Read Strategy	Consistency
MySQL	Single leader	Leader only	Leader / Replica	Strong → Eventual
MongoDB	Primary–Secondary	Primary	Secondary / Primary	Tunable
Cassandra	Leaderless	Quorum	Quorum	Eventual

DB	Replication Model	Write Strategy	Read Strategy	Consistency
DynamoDB	Leaderless	Quorum	Quorum	Eventual / Strong
Redis	Primary–Replica	Primary	Replica	Eventual
Redis Cluster	Partitioned leader	Partition leader	Replica	Eventual

Quick DB Insights

MySQL

- Simple
 - Strong consistency
 - Scaling writes is hard
-

MongoDB

- Flexible read preferences
 - Good balance
 - Still leader-based
-

Cassandra

- Designed for scale
 - Tunable consistency
 - Complex operationally
-

DynamoDB

- Fully managed
 - Global tables
 - Expensive at scale
-

Redis

- Extremely fast
 - Often used as cache
 - Data durability trade-offs
-

ML & Retail Context (Very Important)

Feature Stores

- Async replication
- Read-your-writes required

Online Inference

- Read from nearest replica
- Accept eventual consistency

Training Pipelines

- Replica lag acceptable
 - Throughput more important
-

Architect-Level Summary

- Replication improves availability and read scalability
- Consistency vs availability is always a trade-off
- ML systems must tolerate stale reads
- Leaderless systems scale better but are harder to reason about