

# Lesson 11 -> Database Indexing

## 1 What Is Database Indexing?

A **database index** is a **data structure** that allows **fast lookup of rows** without scanning the entire table.

Think of it like:

- A book index → jump directly to page
  - Instead of reading the entire book
- 

### Without Index

```
SELECT * FROM orders WHERE user_id = 123;
```

- Full table scan
  - $O(N)$
- 

### With Index

```
Index(user_id) → pointer to row
```

- $O(\log N)$
- 

## 2 Why Indexing Is Important

### 1. Faster Reads

- Queries run orders of magnitude faster

### 2. Reduced Disk IO

- Fewer pages read from disk

### 3. Efficient Sorting & Filtering

- ORDER BY
- WHERE
- JOIN

## 4. Enables Scalability

- Without indexes, scaling hardware won't help
- 

## Trade-Off (Important)

Indexes are **not free**:

- Consume disk
  - Slow down writes
  - Need maintenance
- 

## 3 How Indexes Work Internally (Simplified)

Most DBs use:

- B-Tree / B+Tree**
- Some use **LSM Trees** (Cassandra)

Basic idea:

Key → pointer to row location

Example:

user\_id = 123 → page 42, row 7

---

## 4 Types of Indexes

### 4.1 Primary Index (Primary Key Index)

#### What it is

- Index on **primary key**
- Usually **unique**
- Automatically created

#### Characteristics

- Fastest lookups

- Enforces uniqueness

## Example

```
PRIMARY KEY (order_id)
```

## Used for

- Row identity
  - Joins
  - Point lookups
- 

## 4.2 Secondary Index

### What it is

- Index on **non-primary columns**

Example:

```
CREATE INDEX idx_user_id ON orders(user_id);
```

### Characteristics

- Can be non-unique
  - Slower than primary index
  - Essential for filtering
- 

## 4.3 Clustered Index (Very Important)

### What it is

- **Defines physical order of data on disk**
- Table rows stored in index order

### Key Rule

👉 Only one clustered index per table

## Example (MySQL InnoDB)

- Primary key is clustered index

```
Disk: [ PK=1 row ] [ PK=2 row ] [ PK=3 row ]
```

## Pros

- Very fast range queries
- Fewer disk reads

## Cons

- Insert/update expensive
  - Reordering cost
- 

## 4.4 Non-Clustered Index

### What it is

- Index separate from table data
- Stores pointer to row

Index → Row pointer → Data

## Pros

- Multiple per table
- Flexible indexing

## Cons

- Extra lookup (bookmark lookup)
- 

## Clustered vs Non-Clustered (Key Difference)

Aspect	Clustered	Non-Clustered
Physical order	Yes	No
Count per table	One	Many
Speed	Faster for ranges	Slightly slower
Use cases	PK, time-series	Filters, joins

---

## 5 Composite (Multi-Column) Indexes

### What it is

Index on multiple columns.

Example:

```
CREATE INDEX idx_user_date ON orders(user_id, order_date);
```

## Important Rule

- Order matters

Works for:

- user\_id
- user\_id + order\_date

Does NOT work for:

- order\_date alone
- 

## 6 Covering Index

### What it is

- Query can be answered **entirely from index**
- No table lookup needed

Example:

```
SELECT user_id, order_date FROM orders
```

Index contains both columns.

### Benefit

 Extremely fast queries

---

## 7 Indexing in Different Databases

### MySQL (InnoDB)

- Clustered index on primary key
- Secondary indexes point to PK

### PostgreSQL

- No clustered index by default
- Can cluster table manually

## MongoDB

- B-Tree indexes
- Secondary indexes supported

## Cassandra

- LSM-tree based
- Primary key defines partition + clustering

## Redis

- No traditional indexes
  - Uses key-based access
  - Sorted sets for range queries
- 

## 8 Indexing Pitfalls (Architect Thinking)

### Over-Indexing

- Too many indexes → slow writes

### Wrong Column Order

- Composite index misuse

### Low Cardinality Columns

- Index on boolean/status often useless

### Index on Hot Columns

- Causes write contention
- 

## 9 ML & Retail Use Cases

### Feature Store

- Index on entity\_id
- Index on (entity\_id, timestamp)

## **Recommendation Systems**

- Index user\_id
- Index item\_id

## **Order Systems**

- Clustered by order\_date
- Secondary on user\_id